

# Q2 2026 Earnings Call

## Company Participants

- Andy Jassy, President and Chief Executive Officer
- Brian T. Olsavsky, Chief Financial Officer
- Dave Fildes, Vice President of Investor Relations

## Other Participants

- Brian Nowak, Analyst, Morgan Stanley
- Colin Sebastian, Analyst, Baird
- Douglas Anmuth, Analyst, JPMorgan
- Eric Sheridan, Analyst, Goldman Sachs
- Justin Post, Analyst, Bank of America
- Ken Gawrelski, Analyst, Wells Fargo

## Presentation

### Operator

Thank you for standing by. Good day, everyone, and welcome to the Amazon.com Quarter Two 2026 Financial Results Teleconference. At this time, all participants are in a listen-only mode. After the presentation, we will conduct a question-and-answer session. Today's call is being recorded. And for opening remarks, I will be turning the call over to the Vice President of Investor Relations, Mr. Dave Fildes. Thank you, sir. Please go ahead.

### Dave Fildes {BIO 20638976 <GO>}

Hello, and welcome to our Q2 2026 Financial Results Conference Call. Joining us today to answer your questions is Andy Jassy, our CEO; and Brian Olsavsky, our CFO. As you listen to today's conference call, we encourage you to have our press release in front of you, which includes our financial results as well as metrics and commentary on the quarter.

Please note, unless otherwise stated, all comparisons in this call will be against our results for the comparable period of 2025. Our comments and responses to your questions reflect management's views as of today, July 30th, 2026, only and will include forward-looking statements. Actual results may differ materially. Additional information about factors that could potentially impact our financial results is included in today's press release and our filings with the SEC, including our most recent annual report on Form 10-K and subsequent filings.

During this call, we may discuss certain non-GAAP financial measures. In our press release, slides accompanying this webcast, and our filings with the SEC, each of which is posted on our IR website. You will find additional disclosures regarding these non-GAAP measures, including reconciliations of these measures with comparable GAAP measures.

Our guidance incorporates the order trends that we've seen to date and what we believe today to be appropriate assumptions. Our results are inherently unpredictable and may be materially affected by many factors, including fluctuations in foreign exchange rates and energy prices, changes in global economic and geopolitical conditions, tariff in trade policies, resource and supply volatility including from memory chips and customer demand and spending, including the impact of recessionary fears, inflation, interest

rates, regional labor market constraints, world events, the rate of growth of the Internet, online commerce, cloud services, and new and emerging technologies and the various factors detailed in our filings with the SEC.

Our guidance assumes, among other things, that we don't conclude any additional business acquisitions, restructurings, or legal settlements. It's not possible to accurately predict demand for our goods and services, and therefore our actual results could differ materially from our guidance.

And now I'll turn the call over to Andy.

## **Andy Jassy** {BIO 15111610 <GO>}

Thanks, Dave. We're reporting \$200.6 billion in revenue, up 20% year-over-year. Operating income was \$27.5 billion, up 43% year-over-year. Q2 was another very strong quarter for Amazon. I'll start with AWS, which is booming right now, and I'll share the numbers what we think is going on and why we're enthusiastic about the ROIC equation even with heavy CapEx next few years.

First, the numbers. Revenue growth of 36.7% year-over-year, accelerating for the fifth straight quarter, our fastest growth in 18 quarters back when AWS was less than half its current revenue size. We added over \$4.6 billion in revenue quarter-over-quarter, about 80% more than our largest increase ever. And our backlog stands at \$496 billion, growing triple-digits year-over-year. AWS is now a \$169 billion annualized revenue run-rate business, which for perspective would place a 24th on the Fortune 500 list if it was a standalone company. Our chips business now has an annual revenue run rate of over \$25 billion, growing triple-digit percentages year-over-year.

Our AI revenue run-rate climbed significantly quarter-over-quarter and is now also over \$25 billion, growing triple-digit percentages year-over-year. Customers choose AWS because we offer the broadest capabilities they want their AI inference to reside near their other applications and data and more of it resides in AWS than anywhere else and because AWS has the strongest security and operational performance.

We're seeing strong growth across both AI and non-AI, what we call core, and growth in one is driving growth in the other. Growth in AI drives core because post-training reinforcement learning and agent tool use is mostly done on CPUs versus AI accelerators. This is an advantage for AWS as our Graviton chip is the strongest CPU chip, offering up to 30% to 40% better price performance than other options.

You need a place to store this AI data and to run vector databases, which are also emblematic of a meaningful edge for AWS because we have the broadest and most capable functionality by a fair bit in these core infrastructure areas. We feel similarly about the AI stack top-to-bottom.

We have a unique offering that customers are excited about. As we've been saying for 18 months now, technically competent companies are going to build their own foundation models, not the really big frontier models, but smaller models that leverage their proprietary data. There is no easier service for this than our SageMaker AI service.

Customers also need a high-performance, cost-effective inference service, and that's what Amazon Bedrock provides. Bedrock not only provides the best selection of leading

models as superior performance and with the governance of security controls that companies need, it's also continuing to grow incredibly quickly.

In addition to leading model building and inference services, customers need easier ways to build, run and leverage agents. For example, even after you've built an agent, you have a lot of muck to worry about. A production agent needs somewhere secure to run, memory so it holds context, an identity so it can act on a user's behalf, tools and data to connect to and a way to watch what it's doing once real traffic hits.

Stitching all that together reliably is hard, and it solved many production deployments. It's why we've built Bedrock Agent Corp. It (Technical Difficult) managed infrastructure and our teams keep iterating recently adding features like policies, which give companies deterministic controls over what agents can do, payments so agents can execute transactions autonomously, web search to ground agents' knowledge without having to leave AWS, and a new harness that further speeds up how fast customers can put this all together, including creating the agent with strands.

While companies will construct their own purpose-built agents from the ground up, most will also use turnkey agentic services. Coding agents are a good example, and there are several successful ones, including Claude Code, Codex, and our own spec-driven Kiro, which is up to 50% more cost-effective than others and tripled in usage quarter-over-quarter.

Another of these agentic services is Amazon Quick, an intelligent AI work companion that helps you manage, search, and automate your digital workload across email, calendar, local or cloud files, and custom workflows. Unlike other offerings in this space, Quick also lets you manage across leading SaaS tools like Slack, Salesforce, Jira, Teams and ServiceNow.

Quick enforces the company's existing access control, so each person sees only what they're cleared to see, and then it takes action, scheduling meetings, drafting and sending email, updating a CRM record, building a dashboard and more. In Q2, we made Quick even more capable adding autonomous agents to customer setup and playing language to run continuously in the background and carry out multi-step tasks, a personalized activity feed that pulls email, messages, calendars and tasks into one prioritized view and 16 new integrations, including Adobe, Moody's and Snowflake.

Quick has momentum with 3M, Allianz, AstraZeneca, Autodesk, BMW, Exxon, FINRA, Hyundai, Intuit, Mondelez, Moody's, the NBA, the NFL, (Technical Difficulty) Life Insurance, and Southwest Airlines all using it. We also have services like Amazon Connect, our call center service in AWS Transform, which automates software migration growing quickly.

I'll mention one more that I'm excited about. As Frontier models get increasingly powerful, they're making it easier to find security vulnerabilities in technical applications, many of which humans haven't found. This is obviously concerning for companies that protect important data.

We recently released AWS Continuum, which discovers, prioritizes, validates, and remediates code vulnerabilities. It starts by ingesting the backlog of vulnerabilities the team already has and then leverages the new frontier models to run comprehensive scans. Continuum uses agents in each company's own business context to prioritize what matters, reasoning through questions like is the affective component deployed? Is it reachable? Is it in a production path?

And what's the impact if it's exploited? It then validates vulnerabilities in a sandbox, so teams aren't chasing false positives. And finally, it recommends the fix. It's hard to talk with enterprises about AI right now without their mentioning security and we expect continuum to grow quickly.

I mentioned earlier that our chips revenue run-rate is now over \$25 billion. We're unusually well-positioned for this AI inflection given our leading price performance chips in both AI with Trainium and CPU with Graviton. In addition to the two leading AI labs in the world, Anthropic and OpenAI making multi-year, multi-gigawatt commitments to Trainium, an increasing number of AI start-ups are also adopting Trainium, including unicorns like Neurobotics and Odyssey, joining start-ups like 12 Labs, Descarte, Poolside, Karakuri, Metagenomi, Neto Alm and Splash Music and larger companies like Uber and Pinterest all adopting Trainium.

Graviton is used by 98% of our top 1,000 EC2 customers. The revenue commitments have increased nearly three times quarter-over-quarter and Graviton 5 is growing nearly 2x faster as Graviton 4 did. We also continue to have a deep partnership with NVIDIA, and we'll continue making AWS the best place to run NVIDIA chips as we have customers who will run on NVIDIA for as long as we can foresee, and we believe strongly that customers want choice.

Choice is good for customers, competition, and driving the cost of inference down, which customers care deeply about. Let me talk for a second about how we see this investment playing out. Earlier this year, we said we plan to invest approximately \$200 billion in cash CapEx in 2026, the majority of which to support AI and AWS.

At this level of spend and higher, we have clear line-of-sight to strong financial returns. I'll explain why. There are two major parts of the investment, the data centers and the servers and networking equipment that go into them. These have different capital cycles. Data center capital is spent starting two years before we can put servers into them to start monetizing.

Once a data center opens with servers plugged in, we start generating significant revenue right away and then get to monetize these data centers for 30-plus years without having to spend that start-up capital again. Servers and networking equipment operate on a shorter cycle.

We typically purchase these a few months before putting them into service, so we have strong visibility into customer demand before we trigger the spend. If the demand isn't there, we won't spend the capital. For servers and networking equipment, on average, it takes a little less than three years to break even on that investment.

The servers currently have a useful life of at least five to six years and most of our AI capacity these days is being contracted for at least five-year terms. That means that we're driving significant free cash flow on the servers and network of equipment in the two to three years after we break even.

It's also worth noting that AWS has a strong track record of pulling forward breakevens on server equipment, where we've already made meaningful progress and finding ways to extend the useful life of this equipment without sacrificing customer experience. So for our data centers, which have 30-plus-year useful lives, we should get at least five to six generations of server economics, like I explained earlier, with subsequent generations

after the first having even better overall economics because we don't have to repeat that upfront data center investment I mentioned earlier.

This means in the short-term, when demand is necessitating so many data centers being built simultaneously in advance of when we can start monetizing them, we'll spend a lot of CapEx and encounter free cash flow headwinds until these data centers come online, can be monetized and we get a few years into these servers being utilized.

That as we get a few years out and the revenue growth outpaces the incremental CapEx growth, which will happen at some point, the resulting revenue, free cash flow and return on invested capital is very compelling. We've done this before in the first era of cloud computing just over a longer time horizon, where demand built more gradually than it has in AI.

But we see the margins and returns in AI tracking what we saw with core at the same point of evolution, actually a little ahead. We now believe we will spend approximately \$220 billion in cash CapEx in 2026, the higher cost of memory pushing this number up from our prior estimate of about \$200 billion. But even at that amount, we will still not have enough capacity to meet all the demand we have in 2026, and I believe this dynamic will also be true in 2027 too.

In fact, the demand we already have for 2028 is striking. And remember, enterprises are still very early in using inference at scale in their current production applications. We long believed AWS could become a few \$100 billion revenue business, and now believe it will be at least double that and very possibly be \$1 [ph] trillion annual revenue business for us in time with very appealing accompanying free cash flow and return on invested capital.

I'll now turn to stores. We added millions of new products to our selection, including over 700,000 from notable brands. We also expanded ultra-low-price selection on Amazon Hall in the US by nearly 20x since launch and now have over 6 million items priced under \$10. We continue offering everyday low prices that meet or beat our competitors as well as deep discounts of savings during sales events. We're pleased with the customer response to Prime Day, where customers shop millions of deals, including more than 80% at our lowest price of the year and hundreds of thousands discounted by 40% or more.

We're the second largest grocer in the US, and our grocery business continues growing quickly across perishables and non-perishables. The number of monthly active perishables customers grew over 50% since the start of the year. Same-day orders with perishables average over 3x more units per order, and fresh groceries now make up 6 of the top 20 best sellers on Amazon.com.

We grew the number of new customers for Amazon Pharmacy by more than 2x in the first six months of the year and same-day prescription deliveries nearly 5x. We also saved customers nearly \$250 million so far this year in out-of-pocket costs, up more than 400% year-over-year. We continue speeding up overall delivery and once again achieved record delivery speeds for customers in the first half of the year. We offer millions of items for same-day delivery with Prime, up to 40 times more selection than a typical big-box retail store, and our same-day network continues to expand. Globally, we delivered over 40% more items same-day or overnight in the first six months of the year than the same period last year.

We also continued to expand our ultra-fast service, Amazon Now, which offers delivery in 30 minutes or less on thousands of everyday essentials. We added 80 new cities and towns across the US and several major cities in Egypt in Q2, and Amazon Now is available in nine countries in over 250 cities and towns globally. We continue to see strong customer response with over 80% growth in gross sales and units sold quarter over quarter, and we served over 60% more customers quarter over quarter.

We recently launched Amazon Supply Chain Services, so any business can move, store, and deliver everything from raw materials to finished products using the same supply chain that supports Amazon. And we already have several large customers, including Procter & Gamble, 3M, Lands' End, and American Eagle Outfitters. The stores' team also continues to innovate and deliver for customers with AI. Customers love Alexa for shopping or agentic AI shopping assistant. It offers personalized recommendations, product comparisons, price history, and the ability to automate shopping through features like Price Alerts and Auto-Buy.

Over 350 million customers have used it in the last 12 months, and engagement accelerated in Q2 with active users nearly doubling and interactions up over 5x year-over-year. We also expanded Amazon Lens, which lets customers take a photo of anything they see and instantly find the same or similar items on Amazon to 10 additional countries is now available in 21 countries around the world.

Moving on to Amazon ads, we saw strong growth across our offerings, generating \$19.8 billion of revenue, up 26% year-over-year. Sponsored products continues to be our largest offering and a key driver of growth. Additionally, increasingly more shoppers are discovering products in our agentic and conversational experiences, including in Alexa+ and Alexa for Shopping. Shoppers who click a sponsored prompt convert to a sale 48% more often and spent 21% more on average than those who don't. We see continued growth in engagement in Prime Video Ads and live sports. We introduced more than 30 new advertisers to the NBA in our first year, and inventory on Thursday Night Football, NBA, WNBA, and NASCAR all sold out. Advertisers are increasingly investing in multi-sports strategies, with brands activating across multiple sports seeing 2.3 times higher unduplicated reach compared to single-sport advertisers, and multi-sport viewers are driving 12% higher spend and 17% more orders on Amazon.

And finally, we make it easy to create, launch, and optimize full-funnel campaigns using AI-powered tools, including Ads Agent, which turns hours of setup and targeting into minutes. Advertisers using Ads Agent targeting see 8% lower cost-per-impression and 6% lower cost-per-acquisition, and we've expanded it to 11 new countries this year.

We're also continuing to see momentum in several other areas, and I'll mention just a few. Starting with entertainment, the inaugural season of NBA on Prime Video delivered strong viewership with a peak of 6.5 million US viewers for Game 7 of the Eastern Conference Semifinals, outperforming Game 7 on broadcast a year ago. In Europe, viewership of the NBA more than doubled year-over-year on Prime Video, with the highest average viewership on record. We also drew 36 million viewers globally for the series premiere of Off Campus on Prime Video in its first 12 days, becoming Prime Video's number three top-viewed series debut ever.

Alexa+ expanded to Germany, Austria, France, and Brazil, and hundreds of millions of customers are using new Alexa experiences. We find that everywhere Alexa goes, it drives momentum for the business. For example, in the US, customers who use Alexa for

Shopping spend an average of over 40% more per order than those who don't. And customers who've tried Alexa+ are signing up for Prime at a nearly 25% higher rate.

And finally, Amazon Leo has close to 400 satellites in orbit, enough to begin initial satellite internet service this year. We already have meaningful revenue commitments from enterprises and government customers, and we have more than 20 partners who will extend the reach of our network across the globe. We continue to be in the middle of some of the biggest inflections of our lifetime, and we're building multiple new long-term businesses that will make customers' lives better and easier and lead to substantial free cash flow and return on invested capital for our shareholders and business.

With that, I'll turn it over to Brian.

## **Brian T. Olsavsky** {**BIO 18872363 <GO>**}

Thanks, Andy. Starting with our top-line financial results, worldwide revenue was \$200.6 billion, a 20% increase year-over-year, excluding the impact of foreign exchange. These results include a timing shift of Prime Day into Q2 for most of our largest countries, including the United States.

In Q2, we reported worldwide operating income of \$27.5 billion. This includes the benefit from two items that reduced expenses by approximately \$1.2 billion during the quarter. First, we received tariff-related refunds of approximately \$600 million. This is included in our North America segment and represents a significant majority of refunds we expect to receive.

Second, we recorded a separate benefit of approximately \$600 million related to the change in fair-value measurement of energy contracts subject to derivative accounting. This primarily impacts the AWS segment. These energy contracts are to secure electricity supply for existing and future operations. Certain terms in these contracts cause them to be subject to derivative accounting. Derivatives are measured at fair value each reporting period, and changes in fair-value measurements will create unrealized gains and losses recorded within operating expenses on our statements of operations. While the impact of these fair-value measurements can vary, these adjustments have not been significant in prior quarters.

Moving on to our segment results. In the North America segment, second-quarter revenue was \$116.2 billion, an increase of 16% year-over-year. International segment revenue was \$42.2 billion, an increase of 15% year-over-year, excluding the impact of foreign exchange. Worldwide paid units grew 17% year-over-year. Prime remains a key pillar of our business, and our double-digit year-over-year membership growth reflects the value our offerings deliver at scale. We saw broad-based momentum across the inputs that drive our customer experience, including new selection, sharp prices, and fast delivery.

Shifting to profitability, North America segment operating income was \$9.1 billion with an operating margin of 7.9%. International segment operating income was \$1.7 billion with an operating margin of 4.1%. In our fulfillment network, we made progress optimizing inventory placement, shortening shipping distances, reducing touches per package, and improving consolidation rates. We're expanding our deployment of robotics and automation, which have been integral to our operations for decades. We're retrofitting our facilities with our latest generation technology, and we expect to more than double our fleet of robotic arms like Cardinal and Sparrow in 2026. We continue to lower our

overall cost-to-serve even as we face heightened transportation costs driven by fuel inflation from the conflict in the Middle East and higher linehaul rates from driver capacity limitations. Excluding the impact of higher fuel and linehaul rates, shipping costs grew more slowly than worldwide unit growth at a pace that is relatively consistent with last quarter.

On the operating income side, the impact of these higher costs is partially offset by our FBA fuel and logistics surcharge that was implemented in April. Looking ahead, we see meaningful opportunities to further enhance productivity across our global fulfillment network, all while continuing to raise the bar and delivery speed. While operating margin may fluctuate and progress may not always be linear, we take a deliberate approach to achieving sustained long-term improvement in our cost-to-serve. Moving to the AWS segment, revenue was \$42.2 billion, up 36.7% year-over-year, driven by both core and AI services.

AWS now has an annualized revenue run-rate of \$169 billion. Customers continue to increase cloud migrations and scale up their use of AWS core services. And increasingly, customers seeking the full benefits of AI are accelerating their transition to the cloud. We see a strong linkage between AI spend and core growth.

As customers invest in AI, we see a corresponding increase in core consumption. We expect this relationship to strengthen over time as more AI workloads move into full-scale production and drive additional demand for our core services. AWS operating income was \$16.6 billion, which reflects our strong growth coupled with our focus on driving efficiencies across the business.

Our investments in software and process improvements optimized server capacity and help to develop a more efficient network using our lower lower-cost custom silicon and custom network gear. Now turning to our cash CapEx, which is \$53.1 billion in Q2. This primarily relates to AWS and generative AI as we invest to support strong customer demand.

We'll continue to make significant investments, especially in generative AI, as we believe it to be a massive opportunity with the potential to drive long-term revenue and free cash flow. Before moving on to guidance, I'll briefly touch on the impact from the tariff-related refunds on our results.

We are participating in the tariff refund process and as I mentioned earlier, we received approximately \$600 million in Q2. The amount is limited for a couple of reasons. First, our team has did a lot of work forward buying and pre-positioning inventory to avoid tariff costs. Second, we are not the importer of record for the large majority of items sold in our store, given suppliers typically handle imports and pay relevant tariffs.

In cases where we did see an increase in costs due to tariffs, we largely absorbed these costs rather than pass them on to customers. You can see that in how we stayed very sharp in our price throughout the last year with our product prices on average 14% less than other retailers, according to third-party research firm Profitero.

We have identified a limited set of circumstances where we can trace that we pass specific import charges on to customers. And when we receive those refunds, we will proactively contact affected customers and automatically issue refunds to them. Otherwise, like other large retailers, we'll utilize refunds continue to invest in low prices for customers.



I'll finish with our financial guidance. Q3 net sales are expected to be between \$197 billion and \$202 billion. I'll call out two areas that are driving the sequential deceleration in net sales growth from Q2 to Q3. First, Prime Day timing shifted this year with the sales event occurring in Q2 for most of our large countries, including the US.

In 2025, Prime Day was entirely in Q3. Excluding the impact of Prime Day in both 2025 and 2026, third quarter 2026 year-over-year growth would have been nearly 400 basis points higher. We recognize the timing shift of Prime Day creates some noise in the modeling of our financial results.

However, after adjusting for Prime Day, as we look across our stores business heading into Q3, we see strong customer engagement, growth, and a continuation of the trends that we have seen in the first half of the year. Second, the Q3 guidance anticipates an unfavorable impact of approximately 80 basis points from the year-over-year changes in foreign exchange rates based on current rates.

Q3 operating income is expected to be between \$22.5 billion and \$26.5 billion. I want to thank our teams across the company for their hard work and focus on customers. We'll continue to work to deliver more value to the broad range of customers we serve across each of our businesses, which is the only reliable way to create lasting value for our shareholders.

With that, let's move on to your questions.

## Questions And Answers

### Operator

Thank you. At this time, we will now open the call up for questions. (Operator Instructions) And the first question comes from the line of Doug Anmuth with JPMorgan. Please proceed.

### Q - Douglas Anmuth {BIO 5591566 <GO>}

Great. Thanks for taking the questions. Just one for Brian and one for Andy. Brian, many have assumed that AI workloads would be lower margin at least near term. Can you just talk about the drivers of the 39% AWS operating margin in 2Q and just how we should think about sustainability?

And then, Andy, strong Bedrock traction with customers spending more in the quarter than in all the prior quarters combined. But when you think about the full stack offering, does Amazon need its own leading model toward the frontier? Thank you.

### A - Brian T. Olsavsky {BIO 18872363 <GO>}

Doug, let me start with the first question. Yes, we're pleased with the growth in both revenue and also margin expansion that we had in AWS in Q2, especially given the size of our business. You're seeing despite the large investments, AWS margins have continued to remain strong, and we were up 650 basis points year-over-year, 520 basis points if you exclude the derivative accounting gain that I mentioned.

And we said before these margins will fluctuate. They're based on a number of factors, including our investment levels, mix of products, mix of AI versus non-AI. But I would say that the profitability you're seeing from AWS isn't random. It's a result of disciplined

efficiency gains, capacity optimization, which we benefited quite a bit from in Q2, and always closely managing our fixed costs. So again, they'll fluctuate, but very strong performance year-over-year, and we'll take it.

## A - Andy Jassy {BIO 15111610 <GO>}

Yeah. I'll just add one other thing to what Brian said, which is, as I mentioned in my opening comments, we see the AI business following very much the same type of margin trajectory that we saw in the core business before and it's a little bit ahead of that pace that we saw. So we're optimistic about that.

On the question about Bedrock and our own Frontier model, so my view of it is that AWS and Amazon can have a wildly successful business without its own frontier model. And a lot of that is because there is not going to be one model to rule the world. You already see that right now.

You see it. It's not just Anthropic or it's not just OpenAI. You see increasingly more and more companies being interested in the open models as well. And we have all of them in Bedrock. And it's one of the many reasons why Bedrock is growing so quickly. If you're a company that's building important AI applications, you want to make sure that you have the ability to use all the available models.

They're going to each leapfrog each other at different times. They're going to have lots of different models that actually are comparable in capabilities. And you want that leading selection with the right price performance and with the right governance and security, and there's nothing like Bedrock that provides that right now.

And we use those models as well. All that said, we are pursuing our own frontier model and we're doing it for a few reasons. First of which is it just gives us additional control over cost, cost for our own consumer applications, but also we're trying to drive cost down for customers.

And having a player like ourselves, it's always focused on trying to take the price performance and the cost down for customers all the time, we think will help keep the models more cost-effective for customers. I think also it allows us to have more control over prioritization of what models to focus on.

And we have both from our own external customers as well as our internal customers inside the company, certain priorities that matter that we want the models trained especially well for, and then it gives us some control on speed. And so my view of it is that within the next few years you're going to have at least a half-dozen models that are comparably good to each other, they will all be in Bedrock, and one of them will be ours.

## Operator

Thank you. The next question comes from the line of Justin Post with Bank of America. Please proceed with your question.

## Q - Justin Post {BIO 3469195 <GO>}

Great. Thank you. Just thinking about the AWS acceleration, was that really driven by a lot of capacity coming online in the quarter? You guys have been very -- more open than your peers on gigawatts you're adding. Any help you can give us on how much you might

be adding in the second half versus the first half and also how you're thinking about '27? Thank you.

### **A - Andy Jassy {BIO 15111610 <GO>}**

Yeah. I think that the -- there are several reasons for why we saw such significant growth and it's -- we're really pleased and excited fifth straight acceleration quarter, largest acceleration in 18 quarters. And I think there are several things that are driving it. I think the first part is that customers are choosing AWS in part because it has the broadest functionality across both cloud core and AI, in part because it has the strongest operational performance and security and in part because as more and more companies are bringing their inference workloads to production, they want to live near the rest of their workloads and data and so much more of it lives in AWS than anywhere else.

I think there are a couple of other things going on here on the core side. I mean, AI is obviously growing at a very rapid rate, as we talked about, well over \$25 billion in annual revenue run rate at this point. But the core business is growing very quickly as well. And I think there's at least a couple of things at play. One is that increasingly more enterprises are building their transformation plan to move from on-premises to the cloud. Remember, by the way, that 85% of the global IT spend is still on-premises. That equation is going to flip in the next 10 years to 20 years and you see more and more enterprises that are moving and building plans to move to the cloud, and we're winning the lion's share of those with the capabilities I mentioned earlier and the advantages.

And then AI is growing at such a rapid rate and it's pulling along core alongside of it. And that's because the post-training and the reinforcement learning and all the agentic tool use is being driven on CPU and core. And with leading CPU chip-in Graviton, it makes AWS an even more attractive choice. So yes, we're adding a lot of capacity, but there are a lot of other reasons why it's growing. I expect that we will -- we're on pace with the capacity build that we talked about a few quarters ago where we said we expect to have double the capacity -- power capacity by the end of '27 that we had in '25, and we continue to be on that track.

### **Operator**

Thank you. The next question comes from the line of Brian Nowak with Morgan Stanley. Please proceed with your question.

### **Q - Brian Nowak {BIO 16819013 <GO>}**

Thanks for taking my questions. I have two, Andy. I appreciate the color on the long-lived data centers investments versus the server and network investments. Just the question is, as you -- as you sort of look into 2027, you look at the demand that's coming, et cetera, are you at a point where you're going to be able to start to slow that long-lived data center spend at all in '27 or is that just too soon where you're still going to have to be opening up new data centers for the next two, three, four years as you look into '27, that's number one. Second one in the -- in the past 90 days or so, you've talked -- the company has talked publicly about selling Trainium at some point to third-party data centers. How do you think about when you could do that and just sort of the ROIC on that versus core AWS loads?

### **A - Andy Jassy {BIO 15111610 <GO>}**

Well, on the first question, Brian, we have so much demand right now. Apart from what we've talked about in '26, the lion's share of capacity in '27 and we're adding a lot of capacity, as I mentioned just a few minutes ago, is largely reserved. And we have quite a bit of capacity that's already been reserved for '28. And so I think it's actually kind of useful to look at least our view of what we see in the demand and adoption curve right now, which is we see this adoption curve in AI right now is very barbelled. There is on one end of the barbell, the AI labs are consuming gobs and gobs of compute and there are a few runaway successful generative AI applications like Claude Code and ChatGPT.

And on the other end of the barbell are enterprises who are getting real value from AI in cost avoidance and productivity. And these are things like automating customer service or business process automation or fraud or things like that. In the middle of the barbell is all of the current enterprise production workloads, the -- some of which are using inference in a pervasive way, but most of which aren't. And that is going to change very significantly over time. And that will be, in my opinion, that will be the largest absolute segment, the existing production workloads in the enterprise and new businesses and workloads that start-ups build too.

And so I think we're at the -- I think we're still in the relative early stages of how much demand there's going to be for AI. I think it's going to change every customer experience that we know. I think that it will invent all sorts of new ones that we never imagined. And I don't know if the trajectory of that middle part of the barbell will be the same wildly steep trajectory that we've seen with the current barbelled AI labs piece. But we have a lot of demand in front of us and we're going to invest in this business to continue to be the significant market segment leader that we are today. We think, as I mentioned earlier, it has the potential to be \$1 trillion revenue business for AWS, and we intend on continuing to be the leaders.

On the question about selling Trainium. Well, we're quite excited about what's happening in our chips business. As I mentioned earlier, it's over \$25 billion in annual revenue at this point. We think we have the leading price performance chip in both the AI space with Trainium and in the CPU space with Graviton. The fact that we have multiyear, multi-gigawatt commitments from the two largest AI labs in Anthropic and OpenAI, and more and more companies, as I mentioned in my opening comments, using Trainium is exciting and promising. And we have -- we just have an incredible amount of demand for Trainium.

So there are a lot of customers who are very excited about using it in the form that we're providing right now. We do have an increasing number of customers who are interested in us providing the Trainium chips to them separate from even -- from our cloud. And we're actively having those conversations and exploring and I expect there's a real chance we'll do that in the future.

## Operator

Thank you. The next question comes from the line of Colin Sebastian with Baird. Please proceed with your question.

## Q - Colin Sebastian {BIO 6373379 <GO>}

Yeah, good afternoon. Thanks very much. Andy, is it fair to say that there's a more considered effort to move into the application layer with Kiro and Transform and I guess, more broadly plans for workplace productivity tools? And do you see those as sort of providing a boost to the broader platform offering from infrastructure on up? And then,

Brian, I mean, just given the demand signals you guys are both talking about and plans for additional capacity, what are your current thoughts on sources of capital for the build out over the next couple of years? Thank you.

**A - Brian T. Olsavsky {BIO 18872363 <GO>}**

Hi, Colin. Let me take your second question first. You've seen us issue debt this year. We have a lot of options available to us as we continue to fund this growth that we're seeing in AWS. So we'll continue to look at all the options and make the appropriate decision at the right time, but nothing to share today.

**A - Andy Jassy {BIO 15111610 <GO>}**

On your first question, Colin, we see -- we see a very substantial opportunity both for our customers and for AWS in building some of these agentic applications. And some of this is borne out of what customers tell us they wish they had and they want to be using. Some of it is borne out of just needing to provide those capabilities to ourselves inside Amazon. And so Kiro, which is our coding -- agentic coding service is an example of that.

But Amazon Quick is a really interesting example where we just had so many people inside the company who wanted really an intelligent AI assistant to help them work to -- first, it started off with we're a very document-oriented culture. People wanted -- when they got documents not to have to read every document so carefully and to get a summary and then to be able to write their own analyses and responses to these things and to be able to do business intelligence through these agents. And that's really how Amazon Quick started was to do research, to do business intelligence, to do summarization. And then we had so many people inside the company using it that they said, can't you actually find a way to make it much more productive and easier for us to manage our e-mail, to manage our email, to manage our slack communications, to manage our calendar and to use all those things together. And that's really this next instantiation of Amazon Quick. And as I mentioned in my opening comments, it's pretty remarkable not only how fast it's taken off inside Amazon, but how many external enterprises have put it into production with a very large number of people at their companies.

And we see that opportunity up and down kind of the different needs of companies. I think that Connect -- Amazon Connect, which is our call-center service, which is used by all five major leading airline providers as well as many of the leading banks and healthcare companies, continues to grow very quickly. AWS Transform, which makes it much easier to migrate software. It's super useful for enterprises. And then the latest one we just launched with Continuum, it's really hard to have a conversation with a large company about AI right now where they don't actually ask you about security with just all the noise and the and the hype about the security risk with the most current powerful models.

And so Continuum really allows them to use those models productively to find their own vulnerabilities in their code to design the fixes and to help them deploy them. And those are kind of the first set of them. They all have very high promise, but there are several others that we're working on and we think that's going to be very helpful for customers and our business.

**Operator**

Thank you. The next question comes from the line of Jason Helstein with Oppenheimer and Company. Please proceed with your question. Jason, your line is now live. Okay. The next question comes from the line of Ken Gawrelski with Wells Fargo. Please proceed with your question.

**Q - Ken Gawrelski {BIO 16393991 <GO>}**

Thanks so much. Two, if I may. First, your RPO reported is 2.5x that of the third quarter of '25 when you gave us the doubling of capacity comments for year-end '27. How does that RPO number and the massive expansion there impact your outlook for future capacity? I know that you talked about through year-end '27, but maybe at least qualitatively, if you could touch upon what the RPO means for '28 and beyond capacity?

And then second one related is you raised your CapEx guidance this year for some supply-chain inflation. Could you talk about how your pricing strategy at AWS incorporates future cost inflation, your longer-term contracts allow for stable return profiles despite cost inflation? Thank you.

**A - Andy Jassy {BIO 15111610 <GO>}**

Yeah. I'll start on the backlog number. Yeah, it's -- to your point, it's very substantially continuing to grow. I think it's a reflection again of customers being very, very enthusiastic about using AWS, both across core and for AI. We have taken in -- we know about that backlog, obviously. So that's all taken into account in our CapEx projections. And over time, I expect that we will continue to sign more deals with customers. And as I mentioned earlier, we're going to pursue the opportunity to continue being the significant market segment leader that we are.

On the second question on the supply-chain inflation, you know what I would say is that most of the deals that you sign, other than -- there's a certain amount of your demand that is on-demand where there aren't contracts, but a large amount of it tends to be deals and agreements that you've signed. And the deals that you've signed, the -- those will be the prices and those will be the agreements that we have over the duration of that contract. And new agreements that you sign, you always take into account what your costs are in how you ultimately build a price that you agree to with your customers. And I think it's no secret right now to any company in the world that there are inflated prices right now on some of the components like memory and hard drives and SSDs.

**Operator**

Thank you. Our final question comes from the line of Eric Sheridan with Goldman Sachs. Please proceed with your question.

**Q - Eric Sheridan {BIO 22465717 <GO>}**

Thanks so much for taking the question. Maybe pivoting to the commerce business. When you're talking about scaling some of your initiatives around fast commerce and a wider array of supply of groceries and everyday essentials, can you talk a little bit about the signal you're getting from consumers in terms of either adoption rates of those services or what it's doing to overall spend trends and whether there's any countries or geographies where you're seeing different outcomes as you launch and scale some of those services? Thanks so much.



## A - Andy Jassy {BIO 15111610 <GO>}

Yeah. We are -- I'm quite enthusiastic and excited about the pace with which we are growing the amount of everyday essentials as well as perishables in the business right now. And I think some of that has to do with the broader selection we have and some of that has to do with just how fast our delivery has gotten over the last two to three years.

And when you can deliver items to people, when you have that broad selection like we do at low prices and you can deliver selection -- that selection to people as fast as we are right now. People consider you for a lot more of their total purchases and shopping visits. And you know, we've talked for a while about grocery, and we have a very large business in that space.

Last year, it was over \$150 billion in gross merchandise sales, making us the second-largest grocer in the US. And a lot of that are non-perishables. The middle aisles that you'd find in a grocery store where you have consumables and canned goods and beauty items and pharmaceutical items. A big chunk of it is our Whole Foods market business, which is the leading organic grocer out there. And if you look at the growth in geographies that Whole Foods markets are in, they're significantly outpacing the growth of comparable grocers and the profit trajectory continues to trend the right way as well.

And we found a new format there in daily shop in urban settings, it's off to an amazing start that we're expanding very rapidly. And then another -- we always knew that if we wanted to serve the number of customers who want us to serve them and that we want to serve, we had to find a way to offer mass brands and perishables in a significant way.

And we've tried lots of experiments over the last few years as we've talked about on this call for a few years, but we have finally found something that is a real needle mover for us in offering perishables in our same-day facilities that we're now able to offer same-day perishables in 2,300 cities around the US. If you look in those cities, nine of the top 10 best sellers in those geographies are perishables. The number of monthly active perishable customers has increased 50% since the start of this year. And for same-day orders with perishables, on average, they averaged three times more units per order.

And so we're just seeing very significant traction in our everyday essentials and in our grocery items. And we're not done experimenting, by the way with other physical formats on the grocery side, but we've hit on something with same-day perishables in our same-day facilities. It's the changing the trajectory of our Everyday Essentials business.

## Operator

(Technical Difficulty) Joining us on the call today for your questions, a replay will be available on our Investor Relations website for at least three months. We appreciate your interest in Amazon and look forward to speaking with you again next quarter.

---

*This transcript may not be 100 percent accurate and may contain misspellings and other inaccuracies. This transcript is provided "as is", without express or implied warranties of any kind. Bloomberg retains all rights to this transcript and provides it solely for your personal, non-commercial use. Bloomberg, its suppliers and third-party agents shall have no liability for errors in this transcript or for lost profits, losses, or direct, indirect, incidental, consequential, special or punitive damages in connection with the furnishing, performance or use of such transcript. Neither the information nor any opinion expressed in this transcript constitutes a solicitation of*

*the purchase or sale of securities or commodities. Any opinion expressed in the transcript does not necessarily reflect the views of Bloomberg LP. © COPYRIGHT 2026, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.*